

OS-Assignment 2

OS-assignment CSE 3

Github_Link

Submitted to: Sir Amulya Ratna Swain

Author : Dipankar Das

```
From 6b8669f221d9ec34be24b5861e5039ad4994f788 Mon Sep 17 00:00:00 2001
From: Dipankar Das <dipsonu10@hotmail.com>
Date: Sat, 2 Apr 2022 19:58:47 +0530
Subject: [PATCH 3/3] Done Reviewing the code
MIME-Version: 1.0
Content-Type: text/plain; charset=UTF-8
Content-Transfer-Encoding: 8bit
To: amulyafcs@kiit.ac.in
Cc: 20051554@kiit.ac.in
```

LGTM 

Signed-off-by: Dipankar Das <dipsonu10@hotmail.com>

src/proc.c | 7 +-----

1 file changed, 1 insertion(+), 6 deletions(-)

diff --git a/src/proc.c b/src/proc.c

index 97e9dec..69f8c1a 100644

--- a/src/proc.c

+++ b/src/proc.c

@@ -651,20 +651,15 @@ waitpid(int cpid)

 p->state = UNUSED;

 release(&ptable.lock);

 return pid;

-

 }

-

 }

 if(havekids || curproc->killed){

-

 release(&ptable.lock);

 return -1;

-

 }

 sleep(curproc, &ptable.lock);

-

 }

+

 }

 }

--

2.25.1

From 2500744dd14134de885a03ff0304be0227d0420d Mon Sep 17 00:00:00 2001
 From: Dipankar Das <dipsonu10@hotmail.com>
 Date: Sat, 2 Apr 2022 10:43:43 +0530
 Subject: [PATCH 2/3] Added documentation to the scheduling
 MIME-Version: 1.0
 Content-Type: text/plain; charset=UTF-8
 Content-Transfer-Encoding: 8bit
 To: amulyafcs@kiit.ac.in
 Cc: 20051554@kiit.ac.in

[🔒] proc.c

Signed-off-by: Dipankar Das <dipsonu10@hotmail.com>

src/proc.c | 33 ++++++-----
 1 file changed, 21 insertions(+), 12 deletions(-)

diff --git a/src/proc.c b/src/proc.c
 index 9cc4f15..97e9dec 100644

--- a/src/proc.c

+++ b/src/proc.c

@@ -319,17 +319,19 @@ wait(void)

* Scheduler never returns. It loops, doing:

* What it does:

* * there is Time Quantum {Round Robin} --predefined--

- * * choose a process to run from the ready queue which has the

- * highest priority {Priority based} `RUNNABLE`

+ * * choose a process to run from the ready queue which has:

+ * * * the highest priority {Priority based} `RUNNABLE`

+ * * * if same PRI then decide using the Arrival time {FCFS}

* * switch to start running that process

- * * eventually that process transfers control

- * via switch back to the scheduler.

+ * * eventually that process transfers control via switch back to the scheduler.

*

* @how it accomplishes

* as the process are push according to there creation it is already FCFS

* then we select based on the highest priority becomes Priority based

* finally as we are preempting the process according to the Quantum Time

* it is becoming the Round Robin Sched.

+ * TODO:

+ * * to implement Ageing in this scheduling

*/

void

@@ -346,13 +348,19 @@ scheduler(void)

sti();

struct proc *highPri = 0;

```

- // Loop over process table looking for process to run.
  acquire(&ptable.lock);
+ // N = number of process which are 'RUNNABLE'
+ // then  $O(N^2)$  is the algorithm to decide which one to pick
+
  for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
    if(p->state != RUNNABLE)
      continue;

    highPri = p;
+ /**
+  * if one process which is runnable is selected @param highPri
+  * then it take {N} steps to get the best fit process
+  */

    for (p1 = ptable.proc; p1 < &ptable.proc[NPROC]; p1++) {
      if (p1->state != RUNNABLE)
@@ -360,14 +368,16 @@ scheduler(void)

        if (p1->priority < highPri->priority)
          highPri = p1;
-       if (p1->priority == highPri->priority && p1->cr_time < highPri->cr_time)
+
+       // if the Priority based decision is not helping then we are deciding
+       using the
+       // FCFS based method of determination i.e. which one came first
+       if (p1->priority == highPri->priority && p1->cr_time < highPri->cr_time)
          highPri = p1;
    }

    p = highPri;
-   // Switch to chosen process. It is the process's job
-   // to release ptable.lock and then reacquire it
-   // before jumping back to us.
+   // Switch to chosen process. It is the process's job to release ptable.
+   // lock and then reacquire it before jumping back to us.
    c->proc = p;
    switchvm(p);
    p->state = RUNNING;
@@ -375,8 +385,6 @@ scheduler(void)
    swtch(&(c->scheduler), p->context);
    switchkvm();

-   // Process is done running for now.
-   // It should have changed its p->state before coming back.
    c->proc = 0;
  }
  release(&ptable.lock);
@@ -593,7 +601,8 @@ sys_sps(void)
  struct proc *p;
  sti();
  acquire(&ptable.lock);
-  cprintf("R+ -> Running\tR -> Runnable\tS ->
  Sleeping\n\nPID\tNAME\tSTATE\tPRI\tArr_Time\n---\t----\t----\t---\t-----\n");

```

2.25.1

```
[+] foo
[🌀] proc where Round Robin is modified with the Priority Based
referenced NOTES/paper.pdf
[+] Output
```

```

output/Q5.png | Bin 0 -> 81066 bytes
src/Makefile  |    3 +-
src/defs.h    |    2 +-
src/exec.c    |    3 ++
src/foo.c     |   31 ++++++
src/param.h   |    3 +-
src/proc.c    |  100 ++++++-----
src/proc.h    |    3 +-
src/ps.c      |    2 +-
src/syscall.c |    2 +-
src/sysproc.c |    5 ---
11 files changed, 118 insertions(+), 36 deletions(-)
create mode 100644 output/Q5.png
create mode 100644 src/foo.c

```

4 / 11

```

zV_#L&y}CexiTY&$IP!-N#3#}+(&GR9L&VPqXw8iSrZ<sVSmoY`b?WP`>cF)tt~f;g
E3w$@ZuK)15

literal 0
HcmV?d00001

diff --git a/src/Makefile b/src/Makefile
index e4085bf..9824ff6 100644
--- a/src/Makefile
+++ b/src/Makefile
@@ -185,6 +185,7 @@ UPROGS=\
    _zombie\
    _myprocess\
    _cal\
+   _foo\

fs.img: mkfs README $(UPROGS)
    ./mkfs fs.img README $(UPROGS)
@@ -254,7 +255,7 @@ qemu-nox-gdb: fs.img xv6.img .gdbinit
EXTRA=\
    mkfs.c ulib.c user.h cat.c echo.c forktest.c grep.c kill.c\
    ln.c ls.c mkdir.c rm.c stressfs.c usertests.c wc.c prd.c zombie.c\
-   printf.c umalloc.c myprocess.c\
+   printf.c umalloc.c myprocess.c foo.c\
    README dot-bochsrc *.pl toc.* runoff runoff1 runoff.list\
    .gdbinit.tmpl gdbutil\

diff --git a/src/defs.h b/src/defs.h
index d902698..fcaa2a9 100644
--- a/src/defs.h
+++ b/src/defs.h
@@ -122,7 +122,7 @@ void            wakeup(void*);
void            yield(void);
int             getppid(void);
int             sps(void);
-int            waitpid(int);
+int            waitpid(int);
// swtch.S
void            swtch(struct context**, struct context*);

diff --git a/src/exec.c b/src/exec.c
index b40134f..a175e8b 100644
--- a/src/exec.c
+++ b/src/exec.c
@@ -99,6 +99,9 @@ exec(char *path, char **argv)
    curproc->sz = sz;
    curproc->tf->eip = elf.entry; // main
    curproc->tf->esp = sp;
+
+   curproc->priority = 3;
+
    switchvm(curproc);
    freevm(oldpgdir);
    return 0;

```

```

diff --git a/src/foo.c b/src/foo.c
new file mode 100644
index 0000000..f7adb3a
--- /dev/null
+++ b/src/foo.c
@@ -0,0 +1,31 @@
#include "types.h"
#include "stat.h"
#include "user.h"
#include "fcntl.h"
+
+
+int main(int argc, char *argv[]) {
+ int k, n = 3, id;
+ double x = 0, z, d = 1.0;
+
+ id = 0;
+ for (k = 0; k < n; k++) {
+ id = fork();
+ if (id < 0) {
+ //printf(1, "%d failed in fork!\n", getpid());
+ exit();
+ }
+ if (id > 0) {
+ //printf(1, "Parent %d creating child %d\n",getpid(), id);
+ wait();
+ } else {
+ //printf(1, "Child %d created\n",getpid());
+
+ for (z = 0; z < 80000000.00; z += d)
+ x = x + 3.14 * 89.64; // useless calculation
+ break;
+ }
+ }
+
+ exit();
+}
diff --git a/src/param.h b/src/param.h
index a7e90ef..621b741 100644
--- a/src/param.h
+++ b/src/param.h
@@ -1,6 +1,7 @@
#define NPROC      64 // maximum number of processes
#define KSTACKSIZE 4096 // size of per-process kernel stack
-#define NCPU      8 // maximum number of CPUs
+//#define NCPU      8 // maximum number of CPUs
+#define NCPU      1 // maximum number of CPUs
#define NOFILE     16 // open files per process
#define NFILE      100 // open files per system
#define NINODE     50 // maximum number of active i-nodes
diff --git a/src/proc.c b/src/proc.c
index a87573d..9cc4f15 100644
--- a/src/proc.c
+++ b/src/proc.c

```

```

@@ -88,7 +88,7 @@ allocproc(void)
    found:
        p->state = EMBRYO;
        p->pid = nextpid++;
-
+   p->priority = 10;
    release(&ptable.lock);

    // Allocate kernel stack.
@@ -313,18 +313,31 @@ wait(void)
    }
}

-//PAGEBREAK: 42
-// Per-CPU process scheduler.
-// Each CPU calls scheduler() after setting itself up.
-// Scheduler never returns.  It loops, doing:
-//  - choose a process to run
-//  - switch to start running that process
-//  - eventually that process transfers control
-//    via switch back to the scheduler.
+/**
+ * @def scheduler - Per-CPU process scheduler.
+ * Each CPU calls scheduler() after setting itself up.
+ * Scheduler never returns.  It loops, doing:
+ * What it does:
+ * * there is Time Quantum {Round Robin} --predefined--
+ * * choose a process to run from the ready queue which has the
+ *   highest priority {Priority based} `RUNNABLE`
+ * * switch to start running that process
+ * * eventually that process transfers control
+ *   via switch back to the scheduler.
+ *
+ * @how it accomplishes
+ * as the process are push according to there creation it is already FCFS
+ * then we select based on the highest priority becomes Priority based
+ * finally as we are preempting the process according to the Quantum Time
+ *   it is becomming the Round Robin Sched.
+ */
+
void
scheduler(void)
{
    struct proc *p;
+   struct proc *p1;
+
    struct cpu *c = mycpu();
    c->proc = 0;

@@ -332,12 +345,26 @@ scheduler(void)
    // Enable interrupts on this processor.
    sti();

+   struct proc *highPri = 0;

```

```

    // Loop over process table looking for process to run.
    acquire(&ptable.lock);
    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
        if(p->state != RUNNABLE)
            continue;

+       highPri = p;
+
+       for (p1 = ptable.proc; p1 < &ptable.proc[NPROC]; p1++) {
+           if (p1->state != RUNNABLE)
+               continue;
+
+           if (p1->priority < highPri->priority)
+               highPri = p1;
+           if (p1->priority == highPri->priority && p1->cr_time < highPri->cr_time)
+               highPri = p1;
+       }
+
+       p = highPri;
        // Switch to chosen process.  It is the process's job
        // to release ptable.lock and then reacquire it
        // before jumping back to us.
@@ -541,26 +568,49 @@ getppid()
    return myproc()->parent->pid;
}

+// int
+// sys_sps(void)
+// {
+//     struct proc *p;
+//     sti();
+//     acquire(&ptable.lock);
+//     cprintf("PID : PPID : NAME : STATE : CREATE_TIME : SIZE : PRI\n");
+//     for(p = ptable.proc; p<&ptable.proc[NPROC]; p++) {
+//         if(p->state == SLEEPING)
+//             cprintf("%d : %d : %s : SLEEPING : %d : %d : %d\n",p->pid,p->parent-
+>pid,p->name,p->cr_time,p->sz, p->priority);
+//         else if(p->state == RUNNING)
+//             cprintf("%d : %d : %s : RUNNING : %d : %d : %d\n",p->pid,p->parent-
+>pid,p->name,p->cr_time,p->sz, p->priority);
+//         else if (p->state == RUNNABLE)
+//             cprintf("%d : %d : %s : RUNNABLE : %d : %d : %d\n",p->pid,p->parent-
+>pid,p->name,p->cr_time,p->sz, p->priority);
+//     }
+//     release(&ptable.lock);
+//     return 0;
+// }
+
+int
+sys_sps(void)
+{
-     struct proc *p;
-     sti();
-     acquire(&ptable.lock);

```



```

-     cprintf("PID : PPID : NAME : STATE : CREATION TIME : SIZE\n");
-     for(p = ptable.proc; p<&ptable.proc[NPROC]; p++)
-     {
-         if(p->state == SLEEPING)
-             cprintf("%d : %d : %s : SLEEPING : %d : %d\n",p->pid,p->parent->pid,p-
>name,p->cr_time,p->sz);
-         else if(p->state == RUNNING)
-             cprintf("%d : %d : %s : RUNNING : %d : %d\n",p->pid,p->parent->pid,p-
>name,p->cr_time,p->sz);
-         else if (p->state == RUNNABLE)
-             cprintf("%d : %d : %s : RUNNABLE : %d : %d\n",p->pid,p->parent->pid,p-
>name,p->cr_time,p->sz);
-     }
-     release(&ptable.lock);
-     return 0;
-
+ struct proc *p;
+ sti();
+ acquire(&ptable.lock);
+ cprintf("R+ -> Running\tR -> Runnable\tS ->
Sleeping\n\nPID\tNAME\tSTATE\tPRI\tArr_Time\n---\t---\t---\t---\t-----\n");
+ for(p = ptable.proc; p<&ptable.proc[NPROC]; p++) {
+     if (p->pid < 1)
+         continue;
+     char *ch = {'\0'};
+     if (p->state == SLEEPING)
+         ch = "S";
+     else if (p->state == RUNNING)
+         ch = "R+";
+     else if (p->state == RUNNABLE)
+         ch = "R";
+     cprintf("%d\t%s\t%s\t%d\t%d\n", p->pid, p->name, ch, p->priority,p->cr_time);
+ }
+ release(&ptable.lock);
+ return 0;
+ }
+
+
+ int
+ waitpid(int cpid)
+ {
diff --git a/src/proc.h b/src/proc.h
index c1b2d56..13a949e 100644
--- a/src/proc.h
+++ b/src/proc.h
@@ -49,7 +49,8 @@ struct proc {
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;          // Current directory
    char name[16];               // Process name (debugging)
-   int cr_time;                //Process creation time
+   int cr_time;                // Process creation time
+   int priority;               // Process Priority
};

```

```
// Process memory is laid out contiguously, low addresses first:
```

```
diff --git a/src/ps.c b/src/ps.c
```

```
index 21d1679..492662a 100644
```

```
--- a/src/ps.c
```

```
+++ b/src/ps.c
```

```
@@ -5,7 +5,7 @@
```

```
int main(void)
```

```
{
```

```
- printf(1,"20051575\n");
```

```
+ // printf(1,"20051575\n");
```

```
    sps();
```

```
    /*int t = fork();
```

```
    if(t==0){
```

```
diff --git a/src/syscall.c b/src/syscall.c
```

```
index 0f35f40..f829241 100644
```

```
--- a/src/syscall.c
```

```
+++ b/src/syscall.c
```

```
@@ -131,7 +131,7 @@ static int (*syscalls[])(void) = {
```

```
    [SYS_close]    sys_close,
```

```
    [SYS_getppid]  sys_getppid,
```

```
    [SYS_sps]      sys_sps,
```

```
-[SYS_waitpid]    sys_waitpid,
```

```
+ [SYS_waitpid]    sys_waitpid,
```

```
};
```

```
void
```

```
diff --git a/src/sysproc.c b/src/sysproc.c
```

```
index 8a4d3d8..ca252e0 100644
```

```
--- a/src/sysproc.c
```

```
+++ b/src/sysproc.c
```

```
@@ -43,17 +43,12 @@ sys_getpid(void)
```

```
{
```

```
-//-----
```

```
-
```

```
int
```

```
sys_getppid(void)
```

```
{
```

```
    return getpid();
```

```
- //return myproc()->parent->pid;
```

```
}
```

```
-//-----
```

```
-
```

```
int
```

```
sys_sbrk(void)
```

```
{
```

```
--
```

```
2.25.1
```

```
$ foo & ; foo & ; ps ; ps
R+ -> Running    R -> Runnable    S -> Sleeping

PID      NAME      STATE     PRI      Arr_Time
----      -
1         init      S         3         0
2         sh        S         3         3
7         sh        S         10        1988
12        ps        R+        3         1990
9         sh        R         10        1989
11        sh        S         10        1989
R+ -> Running    R -> Runnable    S -> Sleeping

PID      NAME      STATE     PRI      Arr_Time
----      -
1         init      S         3         0
2         sh        S         3         3
7         ps        R+        3         1988
9         sh        R         10        1989
11        sh        S         10        1989
$ ps
R+ -> Running    R -> Runnable    S -> Sleeping

PID      NAME      STATE     PRI      Arr_Time
----      -
1         init      S         3         0
2         sh        S         3         3
18        foo        R         10        2679
17        foo        R         10        2459
9         foo        S         3         1989
11        foo        S         3         1989
16        ps        R+        3         2378
```